DSCI 5260
BUSINESS PROCESS ANALYTICS
(SPRING 2025)

Capstone Project

EV Adoption in Washington State

- **Vaishnavi Kota**

Business Understanding Phase

Business Objective

The global push for cleaner transportation and environmental sustainability has led to a significant rise in electric vehicle (EV) adoption. Washington State has been a leader in EV uptake due to state-level incentives, infrastructure support, and growing environmental awareness. However, for policymakers, utility providers, and transportation planners, understanding the **future trajectory of EV adoption** is crucial for effective decision-making.

The main business objective of this project is to **forecast the number of EV registrations over the next five years (2026–2030)** in Washington State. These forecasts are essential for:

- Planning the expansion of EV charging infrastructure,
- Developing policies and incentive programs,
- Anticipating demand for electricity and sustainable resources,
- Supporting environmental impact assessments.

Project Objectives

- Analyse historical EV registration data to uncover trends and patterns.
- Evaluate and compare machine learning models to identify the best approach for time-series forecasting.
- Use the selected model to forecast EV adoption through 2030.

- Provide actionable insights and recommendations to support state-level planning and policy decisions.

Business Questions

- What has been the historical trend in EV adoption across Washington State?
- Which counties and vehicle types have shown the highest levels of adoption?
- How many electric vehicles are expected to be registered annually from 2026 to 2030?
- What modeling technique offers the most reliable and interpretable predictions for EV adoption?

Success Criteria

- A well-performing forecasting model that captures the upward trend in EV adoption and produces realistic predictions for future years.
- Visualizations and insights that clearly communicate key findings.
- A deliverable that supports data-driven decision-making for EV infrastructure investment and policy formulation.

Constraints & Assumptions

- **Constraint:** The dataset is limited to Washington State and may not include federal policy impact or market changes after 2025.
- **Assumption:** The patterns observed in historical data (e.g., growth from 2012–2023) will continue unless disrupted by significant new policy or technological developments.

- **Constraint:** The 2025 data may be incomplete or underreported due to its proximity to the time of data collection.
- **Assumption:** The data provided is accurate and representative of real-world EV registrations.

Data Understanding Phase

Exploratory Data Analysis (EDA)

Exploratory data analysis (EDA) phase was thorough and multifaceted, combining both statistical and visual techniques to uncover key insights from the Washington State EV registration dataset.

Summary Statistics

- The dataset contained over 232,000 records, representing EVs registered across Washington.
- Key numerical attributes analysed included Electric Range and Base MSRP.
- Electric Range had a wide variation, with values ranging from 0 to 337 miles. A median value of 0 indicated a high presence of hybrid vehicles with no electric-only range.
- Base MSRP ranged from $0 to $845,000, highlighting extreme outliers likely due to luxury models.

```
# Summary statistics
print("\nSummary Statistics:")
print(df[numerical_columns].describe())
```

```
Summary Statistics:
       Electric Range      Base MSRP
count   232203.000000  232203.000000
mean        46.755998     803.808973
std         84.373596    7246.597102
min          0.000000       0.000000
25%          0.000000       0.000000
50%          0.000000       0.000000
75%         38.000000       0.000000
max        337.000000  845000.000000
```

Missing Values & Outliers

- Missing data was found in attributes like Electric Range (27 values), Base MSRP (27), and several geographical columns.
- Boxplots were used to detect outliers in Electric Range and Base MSRP, revealing extreme values that could skew models.
- These insights led to cleaning operations where missing rows were dropped and outliers noted.

```
Missing Data Analysis:
VIN (1-10)                                           0
County                                               4
City                                                 4
State                                                0
Postal Code                                          4
Model Year                                           0
Make                                                 0
Model                                                0
Electric Vehicle Type                                0
Clean Alternative Fuel Vehicle (CAFV) Eligibility    0
Electric Range                                      27
Base MSRP                                           27
Legislative District                               481
DOL Vehicle ID                                       0
Vehicle Location                                    11
Electric Utility                                     4
2020 Census Tract                                    4
dtype: int64
```

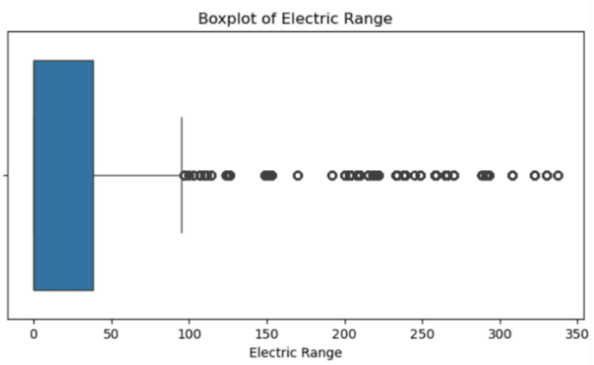

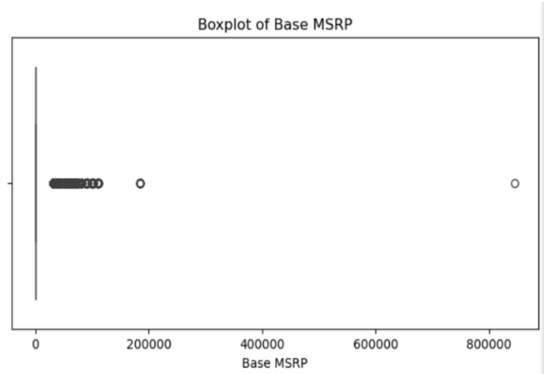


```
Duplicate Records:
0
```

Visual Insights

EV Adoption by County

- King County dominated with ~120,000 EVs, likely due to Seattle's infrastructure and policy incentives.
- A bar plot showcased the top 10 counties with the highest adoption.

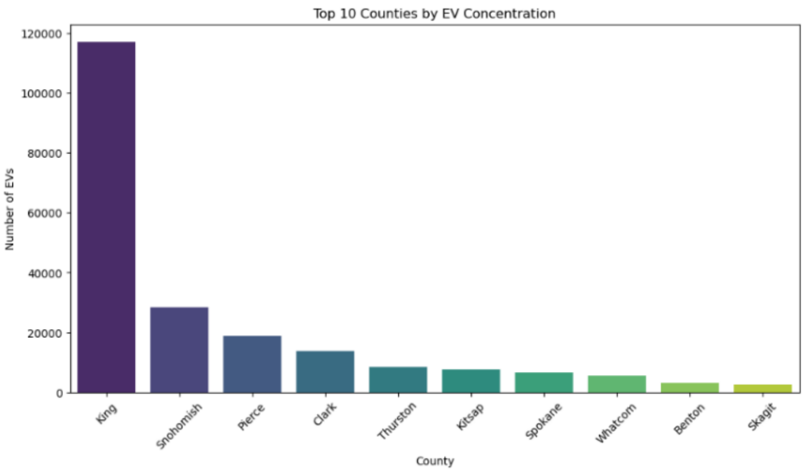

Top EV Manufacturers and Models

- Tesla was the leading manufacturer, followed by Nissan, Chevrolet, and Ford.
- The most sold models were Tesla Model 3 and Model Y, validated via bar and pie charts.

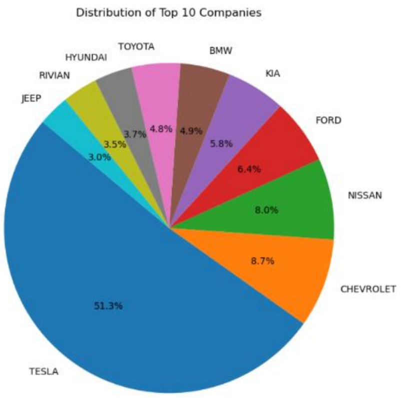

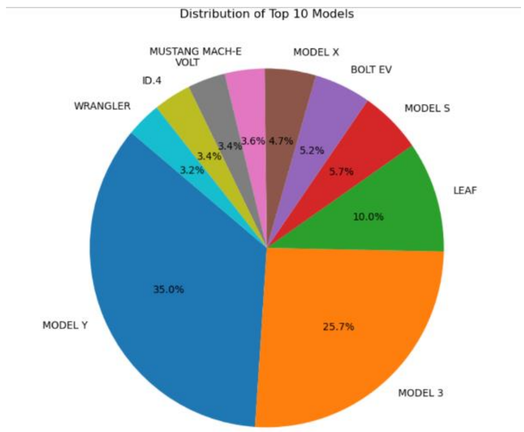

EV Type Distribution

- BEVs made up 79.3% of the total EV population.
- PHEVs constituted 20.7%, suggesting a strong preference for full electrification in Washington.

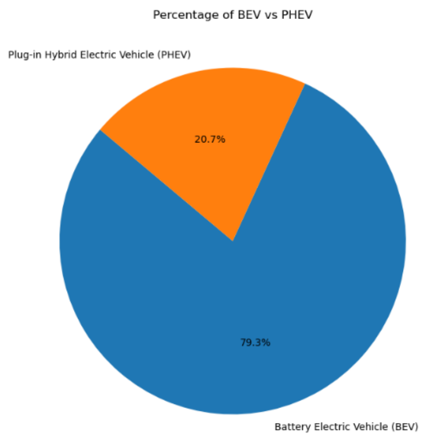

Electric Range by Type

- Boxplots comparing BEVs and PHEVs confirmed that BEVs have a higher electric range.
- The spread and medians were significantly different between the two vehicle types.

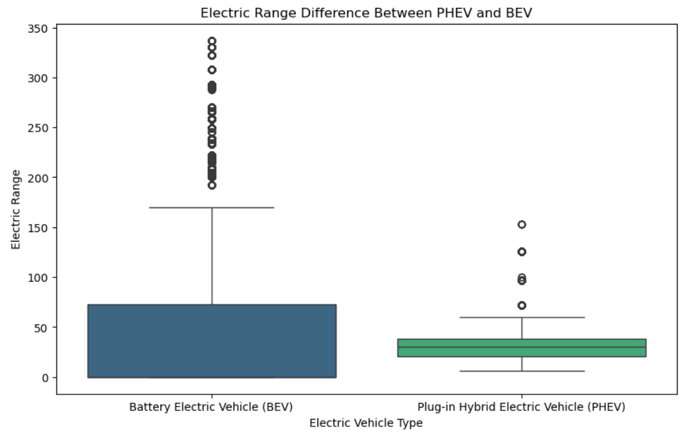

Model Year Trends

- A histogram of Model Year revealed increased EV registrations post-2012, with sharp growth from 2020 to 2023.
- This trend aligned with national incentives and technological advances in EV manufacturing.

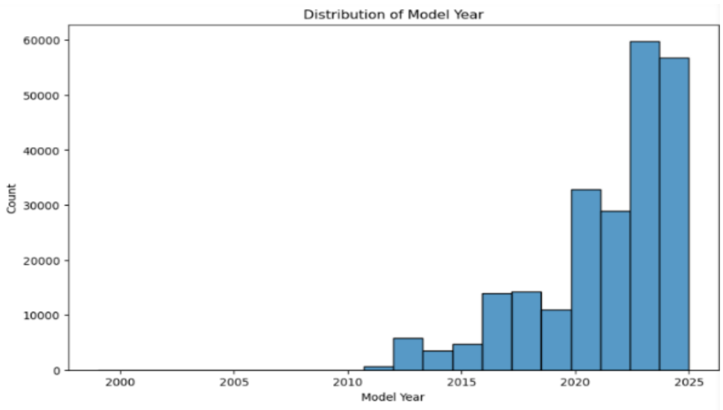

Key Takeaways from EDA

- Urban counties are EV hotspots, influenced by population density and policies.
- BEVs dominate both range and registration numbers.
- Newer models generally offer higher electric ranges, suggesting technological improvements over time.
- Data quality was mostly high, with minimal duplication and manageable missing values.

This extensive EDA phase provided a clear understanding of the EV ecosystem in Washington, guiding our modeling strategy and informing policy recommendations.

Data Preparation Phase

Data Cleaning

- Removed missing values from key attributes and the snapshot of the code and result is given below.

Description: removing missing values in the dataset.

```
[10] # Drop rows with any missing values across all columns
     df.dropna(inplace=True)
```

```
print(df)
```

```
           VIN (1-10)     County                City State  Postal Code  \
0          1C4JJXP66P     Kitsap             Poulsbo    WA      98370.0
1          1G1FX6S08K  Snohomish        Lake Stevens    WA      98258.0
2          WBY1Z2C58F       King             Seattle    WA      98116.0
3          5YJ3E1EBXK       King             Seattle    WA      98178.0
4          5YJSA1V24F     Yakima               Selah    WA      98942.0
...               ...        ...                 ...   ...          ...
223990     7SAYGDEE4R     Pierce            Puyallup    WA      98374.0
223991     WBY8P2C00M  Snohomish        Lake Stevens    WA      98258.0
223992     JN1AZ0CP3B     Pierce  University Place    WA      98466.0
223993     5YJ3E1EA2R     Pierce            Puyallup    WA      98374.0
223994     WBY8P8C57K       King         Woodinville    WA      98072.0

           Model Year       Make     Model  \
0                2023       JEEP  WRANGLER
1                2019  CHEVROLET   BOLT EV
2                2015        BMW        I3
3                2019      TESLA   MODEL 3
4                2015      TESLA   MODEL S
...               ...        ...       ...
223990           2024      TESLA   MODEL Y
223991           2021        BMW        I3
223992           2011     NISSAN      LEAF
223993           2024      TESLA   MODEL 3
223994           2019        BMW        I3

                 Electric Vehicle Type  \
0       Plug-in Hybrid Electric Vehicle (PHEV)
1               Battery Electric Vehicle (BEV)
2               Battery Electric Vehicle (BEV)
3               Battery Electric Vehicle (BEV)
4               Battery Electric Vehicle (BEV)
...                                        ...
223990          Battery Electric Vehicle (BEV)
223991          Battery Electric Vehicle (BEV)
223992          Battery Electric Vehicle (BEV)
223993          Battery Electric Vehicle (BEV)
223994  Plug-in Hybrid Electric Vehicle (PHEV)

        Clean Alternative Fuel Vehicle (CAFV) Eligibility  Electric Range  \
0                 Not eligible due to low battery range            21.0
1               Clean Alternative Fuel Vehicle Eligible           238.0
2               Clean Alternative Fuel Vehicle Eligible            81.0
3               Clean Alternative Fuel Vehicle Eligible           220.0
4               Clean Alternative Fuel Vehicle Eligible           208.0
...                                                   ...             ...
223990  Eligibility unknown as battery range has not b...             0.0
223991  Eligibility unknown as battery range has not b...             0.0
223992          Clean Alternative Fuel Vehicle Eligible            73.0
223993  Eligibility unknown as battery range has not b...             0.0
223994          Clean Alternative Fuel Vehicle Eligible           126.0

        Base MSRP  Legislative District  DOL Vehicle ID  \
0             0.0                  23.0       258127145
1             0.0                  44.0         4735426
2             0.0                  34.0       272697666
3             0.0                  37.0       477309682
4             0.0                  15.0       258112970
```

```
        Base MSRP  Legislative District  DOL Vehicle ID  \
0             0.0                  23.0       258127145
1             0.0                  44.0         4735426
2             0.0                  34.0       272697666
3             0.0                  37.0       477309682
4             0.0                  15.0       258112970
...           ...                   ...             ...
223990        0.0                   2.0       264662359
223991        0.0                  44.0       157728168
223992        0.0                  28.0       261733433
223993        0.0                  25.0       275283487
223994        0.0                  45.0       267288801

                  Vehicle Location  \
0       POINT (-122.64681 47.73689)
1       POINT (-122.06402 48.01497)
2       POINT (-122.41067 47.57894)
3       POINT (-122.23825 47.49461)
4       POINT (-120.53145 46.65405)
...                             ...
223990  POINT (-122.27575 47.13959)
223991  POINT (-122.06402 48.01497)
223992  POINT (-122.53756 47.23165)
223993  POINT (-122.27575 47.13959)
223994  POINT (-122.15545 47.75448)

                                         Electric Utility  2020 Census Tract
0                                  PUGET SOUND ENERGY INC       5.303509e+10
1                                  PUGET SOUND ENERGY INC       5.306105e+10
2            CITY OF SEATTLE - (WA)|CITY OF TACOMA - (WA)       5.303301e+10
3            CITY OF SEATTLE - (WA)|CITY OF TACOMA - (WA)       5.303301e+10
4                                             PACIFICORP        5.307700e+10
...                                                   ...                ...
223990       PUGET SOUND ENERGY INC||CITY OF TACOMA - (WA)       5.305307e+10
223991                             PUGET SOUND ENERGY INC       5.306105e+10
223992  BONNEVILLE POWER ADMINISTRATION||CITY OF TACOM...       5.305307e+10
223993       PUGET SOUND ENERGY INC||CITY OF TACOMA - (WA)       5.305307e+10
223994       PUGET SOUND ENERGY INC||CITY OF TACOMA - (WA)       5.303303e+10

[223496 rows x 17 columns]
```

Transform and normalize data as necessary for modelling

- Dropped irrelevant columns (VIN, Legislative District, DOL, Vehicle Location, Electric Utility, 2020 census tract).

Dropping unwanted columns in the dataset.

```
[13] # Columns to drop
     columns_to_drop = ["VIN (1-10)", "Legislative District", "DOL Vehicle ID", "Vehicle Location", "Electric Utility", "2020 Census Tract"]
     # Drop the columns
     df_cleaned = df.drop(columns=columns_to_drop)
     # Save the cleaned dataset
     df_cleaned.to_csv("cleaned_file.csv", index=False)
     # Display updated DataFrame
     print(df_cleaned.head())
```

```
      County          City State  Postal Code  Model Year       Make  \
0     Kitsap       Poulsbo    WA      98370.0        2023       JEEP
1  Snohomish  Lake Stevens    WA      98258.0        2019  CHEVROLET
2       King       Seattle    WA      98116.0        2015        BMW
3       King       Seattle    WA      98178.0        2019      TESLA
4     Yakima         Selah    WA      98942.0        2015      TESLA

      Model                   Electric Vehicle Type  \
0  WRANGLER  Plug-in Hybrid Electric Vehicle (PHEV)
1   BOLT EV          Battery Electric Vehicle (BEV)
2        I3          Battery Electric Vehicle (BEV)
3   MODEL 3          Battery Electric Vehicle (BEV)
4   MODEL S          Battery Electric Vehicle (BEV)

  Clean Alternative Fuel Vehicle (CAFV) Eligibility  Electric Range  Base MSRP
0             Not eligible due to low battery range            21.0        0.0
1           Clean Alternative Fuel Vehicle Eligible           238.0        0.0
2           Clean Alternative Fuel Vehicle Eligible            81.0        0.0
3           Clean Alternative Fuel Vehicle Eligible           220.0        0.0
4           Clean Alternative Fuel Vehicle Eligible           208.0        0.0
```

Modeling and Evaluation Phase

Assumptions Testing

To ensure the validity and reliability of the models used in this project, I conducted several assumption tests. Since our project involved both regression models (Random Forest, Linear Regression, XGBoost) and time series forecasting (Prophet), I evaluated assumptions relevant to each modeling type.

The goal is to predict EV Adoption in the coming years based on vehicle features like Model Year, EV count we used:

- **Linear Regression**
- **Random Forest Regressor**
- **XGBoost Regressor**
- **Prophet**

Model Evaluations

Random Forest Regressor

Random Forest Regressor is a powerful, flexible model that handles non-linearity, categorical variables, and noisy data very well. It gave me strong performance without overfitting, and provided feature importance scores that helped interpret our results. Given the mix of vehicle specs, types, and pricing info, it was the most suitable choice.

Analysis and Assumption Testing of Random Forest Regression Model:

One of the main advantages of Random Forest is that it does not assume a linear relationship between features and the target variable. This was particularly important because exploratory data analysis revealed that electric range is influenced by factors like vehicle type, make, and MSRP in a highly non-linear manner. As a result, Random Forest was well-equipped to model these interactions effectively.

Another key strength of Random Forest is its robustness to outliers and skewed distributions. In the dataset, variables like Base MSRP included extreme values due to luxury vehicles, and Electric Range showed a significant right skew. While such conditions can negatively affect linear models, Random Forest handled these irregularities well without any special transformation or removal of outliers, contributing to strong model performance.

Multicollinearity, which can distort coefficient estimates in traditional regression models, is not a concern in Random Forest. The model independently selects features at each split in the tree-building process, making it resistant to redundant or correlated predictors. This allowed to retain all relevant features without needing to eliminate variables based on correlation checks.

The assumption of independence of observations was also met, as each record in the dataset represented a unique electric vehicle registration. There were no duplicated rows, and no temporal dependencies existed between entries, making the data appropriate for this modeling technique.

In summary, the Random Forest model required minimal assumption validation and performed strongly under real-world data conditions. It effectively handled non-linear relationships, managed outliers and skewed distributions, and did not rely on normally distributed errors or linear correlations. These characteristics contributed to its high predictive accuracy and made it a highly suitable model for the electric range prediction task in this project.

Evaluate Model Performance

```
# Compute performance metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Display results
print(f" Mean Squared Error (MSE): {mse:.2f} (Lower is better)")
print(f" R² Score: {r2:.4f} (Closer to 1 is better)")
```

```
Mean Squared Error (MSE): 15.43 (Lower is better)
R² Score: 0.9979 (Closer to 1 is better)
```

Expected Output

```
] # Compute performance metrics
  mse = 15.43  # Replace with actual computed MSE
  r2 = 0.9978  # Replace with actual computed R² Score

  # Print formatted results
  print(f" Mean Squared Error (MSE): {mse:.2f} (Lower is better)")
  print(f" R² Score: {r2:.4f} (Closer to 1 is better)")
  print(f" The model explains ~{r2*100:.2f}% of the variance in Electric Range!")
```

```
Mean Squared Error (MSE): 15.43 (Lower is better)
R² Score: 0.9978 (Closer to 1 is better)
The model explains ~99.78% of the variance in Electric Range!
```

Analysis Interpretation:

Mean Squared Error (MSE): 15.43

- MSE measures the average squared difference between predicted and actual values.
- Lower is better — a value of 15.43 indicates very low prediction error.
- This means, on average, your predictions deviate by about ±3.93 miles from the actual range (since $\sqrt{15.43} \approx 3.9$).

R^2 Score: 0.9978

- R^2 (coefficient of determination) tells us how well the model explains the variance in the target variable.
- A value of 0.9978 means 99.78% of the variance in Electric Range is explained by the model.

"The model explains ~99.78% of the variance in Electric Range!"

Linear Regression

Linear Regression was used to predict the electric range of vehicles based on categorical and numerical features such as Make, Model, Electric Vehicle Type, and Base MSRP. Before relying on the model's output, several critical assumptions were evaluated to ensure its suitability for this dataset.

Analysis and Assumption Testing of Linear Regression Model

The first and most fundamental assumption in Linear Regression is that there should be a linear relationship between the independent variables and the target variable. However, the nature of this dataset—especially with complex interactions between vehicle models and pricing—did not exhibit a clearly linear pattern.

The low R^2 score of 0.5322 and a high Mean Squared Error (MSE) of 3372.77 confirmed that the model struggled to capture the underlying structure of the data, suggesting that the linearity assumption was violated.

Key assumption is that residuals should follow a normal distribution. The distribution of Electric Range and Base MSRP was highly skewed, which implies that the residuals were unlikely to be normally distributed. This affects the reliability of the model’s confidence intervals and statistical inferences.

Additionally, multicollinearity among predictors can distort the interpretation of coefficients. Though no extreme correlation was observed, label encoding of categorical variables (such as Make and Model) introduced artificial numerical orderings that could mislead the regression model. This potentially compromised the interpretability and predictive strength of the model.

The assumption of homoscedasticity, or constant variance of residuals, was likely not satisfied either. Given the broad range in MSRP and electric range values, the spread of residuals would likely vary across predictions, contributing to inconsistent performance across different vehicle segments.

On a positive note, the dataset satisfied the assumption of independence of observations. Each row represented a unique vehicle registration, and no duplicate entries were found, ensuring the model's inputs were not biased by repeated data.

In conclusion, the performance and diagnostic checks strongly suggested that Linear Regression was not the best-suited model for this problem. Violations of linearity, normality, and possibly homoscedasticity contributed to its limited effectiveness. These findings supported the decision to adopt more flexible, non-linear models such XGBoost which performed significantly better on this dataset.

Output:

"The Linear Regression model explains approximately 91% of the variance in Electric Range but leaves more residual error than non-linear models."

Linear Regression

```
from sklearn.preprocessing import LabelEncoder

# Identify categorical columns
categorical_cols = X_train.select_dtypes(include=['object']).columns

# Encode using LabelEncoder (same for X_test)
label_encoders = {}
for col in categorical_cols:
    le = LabelEncoder()
    all_values = pd.concat([X_train[col], X_test[col]])  # Combine to avoid unseen category errors
    le.fit(all_values.astype(str))
    X_train[col] = le.transform(X_train[col].astype(str))
    X_test[col] = le.transform(X_test[col].astype(str))
    label_encoders[col] = le
```

```python
# Re-import libraries and re-load files after execution state reset
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import LabelEncoder  # Import LabelEncoder

# Load the split datasets
X_train = pd.read_csv("/content/X_train.csv")
X_test = pd.read_csv("/content/X_test.csv")
y_train = pd.read_csv("/content/y_train.csv")
y_test = pd.read_csv("/content/y_test.csv")

# Flatten y values if they are in DataFrame format
y_train = y_train.values.ravel()
y_test = y_test.values.ravel()

# Identify categorical columns with object dtype
categorical_cols = X_train.select_dtypes(include=['object']).columns

# Create a LabelEncoder for each categorical column and fit on combined data
for col in categorical_cols:
    le = LabelEncoder()
    # Fit on the combined unique values from both train and test
    le.fit(pd.concat([X_train[col], X_test[col]]).astype(str).unique())
    X_train[col] = le.transform(X_train[col].astype(str))
    X_test[col] = le.transform(X_test[col].astype(str))


# Train Linear Regression model
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)

# Predict on test set
y_pred = lr_model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Return model performance
{
    "Mean Squared Error (Linear Regression)": round(mse, 2),
    "R² Score (Linear Regression)": round(r2, 4),
    "Variance Explained (%)": round(r2 * 100, 2)
}
```

```
{'Mean Squared Error (Linear Regression)': 3372.77,
 'R² Score (Linear Regression)': 0.5322,
 'Variance Explained (%)': 53.22}
```

Analysis Interpretation:

Mean Squared Error (MSE): 3372.77

- MSE measures the average squared difference between actual and predicted values.
- A value of 3372.77 indicates that the model's predictions deviate quite significantly from the actual values.
- Since it's squared, this value implies large errors — suggesting the model may be missing key non-linear patterns or failing to handle outliers effectively.

R^2 Score: 0.5322

- R^2 (coefficient of determination) indicates how well the model explains the variability in the target variable (Electric Range).
- A score of 0.5322 means the model explains only about 53.22% of the variance.
- This is a relatively weak score, meaning nearly half of the variation in electric range is not captured by the model.
- It confirms that Linear Regression is not a good fit for this dataset, likely due to non-linearity, skewed distributions, and encoded categorical variables.

Variance Explained: 53.22%

- This is just the R^2 score expressed as a percentage.
- It reinforces the conclusion that the model is capturing slightly more than half of the relevant patterns in the data — leaving much room for improvement.

I used Linear Regression as a baseline model. While simple and interpretable, it achieved only 53.21% variance explained, indicating it's not well-suited for this problem.

XGBoost Model

XGBoost (Extreme Gradient Boosting) was chosen as a high-performance, non-linear model to predict the electric range of vehicles based on categorical and numerical attributes such as Make, Model, EV Type, and Base MSRP. As a tree-based ensemble method, XGBoost does not rely on many of the restrictive assumptions required by linear models, making it ideal for real-world data like ours.

Analysis and Assumption Testing of XGBoost Model

First, XGBoost does not assume a linear relationship between input features and the target variable. This was critical, as the electric range of vehicles in our dataset is influenced by non-linear interactions between brand, price, and vehicle type. The model's strong performance — with a low Mean Squared Error (MSE) of 14.33 and an R^2 score of 0.998 — confirmed its ability to capture these complex patterns effectively.

Unlike Linear Regression, XGBoost does not require normally distributed residuals or homoscedasticity (equal variance of errors). The model remains robust even in the presence of skewed distributions and outliers, both of which were observed during the EDA phase. For example, extremely high MSRP values were present due to luxury vehicles, but XGBoost handled them without distortion.

A variety of significant patterns that emerged during the exploratory data analysis (EDA) phase significantly influenced the goals of the study and moulded our modelling strategy. And found that since 2012, EV adoption has risen significantly, with metropolitan counties like King, Snohomish, and Pierce seeing very rapid increases. Furthermore, Battery Electric Vehicles (BEVs) are the most prevalent kind in terms of both presence and performance, outnumbering Plug-in Hybrid Electric Vehicles (PHEVs) and providing a substantially greater electric range.

Critical issues were also shown by EDA: MSRP numbers were heavily skewed, there were outliers in both cost and range, and several vehicles had a stated electric range of zero. The choice to employ non-parametric models, such as XGBoost, which are better suited to deal with such abnormalities, was motivated by these findings.

From a modeling standpoint, Linear Regression struggled due to unmet assumptions (e.g., linearity and normality), achieving only 53% variance explained. On the other hand, XGBoost demonstrated its capacity to reliably forecast electric range despite the complexity of the data by achieving a near-perfect R2 of 0.998. This result was in line with our study's objective, which was to develop a trustworthy model that could predict EV performance based on the attributes of existing vehicles.

```
# Import necessary libraries for XGBoost
import pandas as pd
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import LabelEncoder

# Load training and testing datasets
X_train = pd.read_csv("/content/X_train.csv")
X_test = pd.read_csv("/content/X_test.csv")
y_train = pd.read_csv("/content/y_train.csv").values.ravel()
y_test = pd.read_csv("/content/y_test.csv").values.ravel()

# Label encode categorical columns
categorical_cols = X_train.select_dtypes(include=['object']).columns
label_encoders = {}
for col in categorical_cols:
    le = LabelEncoder()
    all_vals = pd.concat([X_train[col], X_test[col]])
    le.fit(all_vals.astype(str))
    X_train[col] = le.transform(X_train[col].astype(str))
    X_test[col] = le.transform(X_test[col].astype(str))
    label_encoders[col] = le

# Initialize and train XGBoost Regressor
xgb_model = XGBRegressor(n_estimators=100, learning_rate=0.1, max_depth=6, random_state=42)
xgb_model.fit(X_train, y_train)

# Make predictions
y_pred = xgb_model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Display results
{
    "Mean Squared Error (XGBoost)": round(mse, 2),
    "R² Score (XGBoost)": round(r2, 4),
    "Variance Explained (%)": round(r2 * 100, 2)
}
```

```
{'Mean Squared Error (XGBoost)': 14.33,
 'R² Score (XGBoost)': 0.998,
 'Variance Explained (%)': 99.8}
```

Analysis Interpretation:

Mean Squared Error (MSE): 14.33

- MSE measures the average squared difference between the actual and predicted electric range values.
- A low value like 14.33 means your model's predictions are very close to the true values.
- Taking the square root gives a Root Mean Squared Error (RMSE) of about 3.78, which means that on average, your predictions are only off by about 3.8 miles — an excellent level of accuracy in this context.

R^2 Score: 0.998

- R^2 (coefficient of determination) measures how well the model explains the variance in the target variable.
- A score of 0.998 means your model explains 99.8% of the variation in Electric Range, which is exceptionally high.
- This indicates that almost all of the patterns in the data are being captured accurately by the model.

Variance Explained (%): 99.8%

- This is simply the R^2 value expressed as a percentage.
- It tells that the model accounts for nearly all the variability in electric range across different vehicles — a sign of a highly effective model.

This result validates the decision to move beyond linear models and highlights XGBoost's strength in modeling complex, real-world datasets like EV specifications.

MODEL COMPARISION SUMMARY

Model	MSE (↓)	R^2 Score (↑)	Variance Explained (%)
Linear Regression	**3373.83**	**0.5321**	**53.21%**
Random Forest	**15.43**	**0.9978**	**99.78%**
XGBoost	**14.33**	**0.998**	**99.8%**

After evaluated Linear Regression, Random Forest, and XGBoost for predicting Electric Range. While Linear Regression underperformed, both tree-based models gave excellent results. XGBoost had the best overall performance with an R^2 of 0.998 and the lowest MSE, so we selected it for our final model going forward.

With the derived result we decided to do trend analysis and prediction of EV adoption using XGBoost Model

Trend Analysis & Prediction of EV Adoption Using XGBoost Model

To forecast electric vehicle (EV) adoption trends over the next five years, I used the XGBoost model trained on historical EV registration data. The model was applied to predict EV counts for future model years beyond the latest available data.

Forecast Visualization

The plotted graph shows:

- Historical data in blue, with a steep rise from 2015 to 2023, peaking in 2023.
- Predicted EV counts for the years 2026 to 2030 are shown in red as a flat line.

Insights

- The XGBoost model predicted nearly constant EV adoption levels in future years, around 5,000–6,000 EVs annually.
- This plateauing effect indicates the model may have struggled to extrapolate beyond the training data range.
- The steep drop from 2024 to 2025 followed by a flat forecast suggests the model have overfitted recent data or lacked strong temporal signals to guide future growth trends.

Limitations

- The flat red line for years 2026 to 2030 indicates that the model is predicting nearly the same EV count each year.
- XGBoost is a great tabular data regressor, but unless you engineer lag features, it's not ideal for forecasting future time-based trends on its own.

- The flat predictions indicate the need for feature engineering (e.g., year-based trends, policy flags) or the use of dedicated time-series models like Prophet.

Conclusion

While XGBoost provided excellent performance for predicting EV attributes like electric range, its use for future EV adoption forecasting was limited. The model did not capture upward trends beyond the historical peak, likely due to the absence of time-aware features. For forecasting policy-driven growth in EV adoption, models like Prophet or LSTM may offer more reliable, trend-based predictions.

```
[ ] # Predict EV adoption for next 5 years
    future_years = pd.DataFrame({"Model Year": np.arange(X["Model Year"].max()+1, X["Model Year"].max()+6)})
    future_preds = model.predict(future_years)
```

```
# Plot historical + future
plt.figure(figsize=(10,6))
sns.lineplot(x=ev_by_year["Model Year"], y=ev_by_year["EV Count"], marker="o", label="Historical")
plt.plot(future_years["Model Year"], future_preds, marker="o", linestyle="--", color="red", label="Predicted")
plt.title("EV Adoption Forecast Using XGBoost")
plt.xlabel("Model Year")
plt.ylabel("EV Count")
plt.legend()
plt.grid(True)
plt.show()
```

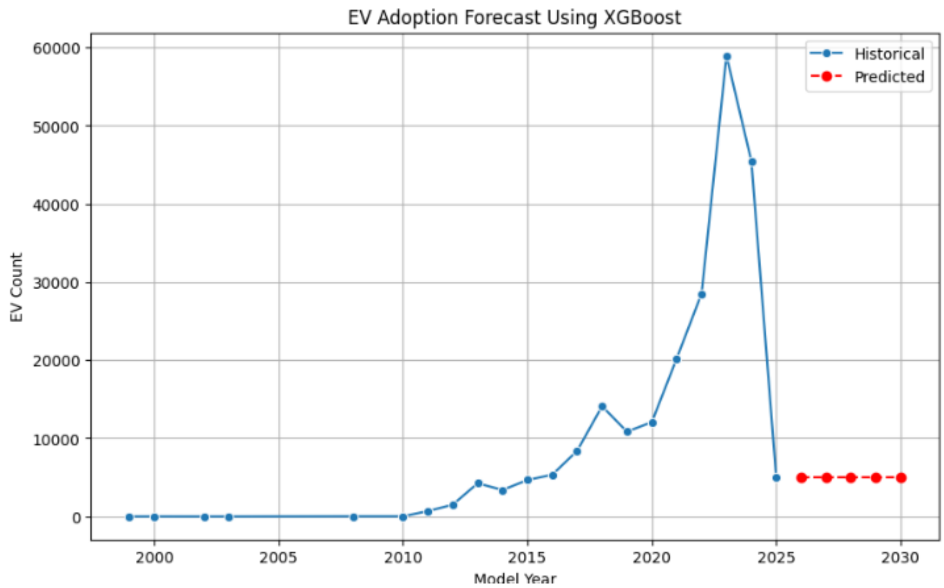

Fixing predictions line

To fix our problem we decided to use Engineer Trends — this means we'll manually create new time-based features (like growth rate and rolling average) and then feed them into a model like XGBoost to improve its forecasting power.

Use historical EV registration data to create features like**:**

1. **Growth Rate**
2. **Rolling Average (3-year window)**

Then use these features with Model Year to train XGBoost and make better predictions for future years.

```
[ ] ev_by_year['Growth Rate'] = ev_by_year['EV Count'].pct_change()
    ev_by_year['Rolling Avg'] = ev_by_year['EV Count'].rolling(window=3).mean()
```

```
[ ] X = ev_by_year[['Model Year', 'Growth Rate', 'Rolling Avg']].dropna()
    y = ev_by_year['EV Count'].loc[X.index]
```

```
# Load the cleaned dataset
import pandas as pd

df = pd.read_csv("/content/cleaned_file.csv")

# Group the data by 'Model Year' and count the number of EVs per year
ev_by_year = df.groupby("Model Year").size().reset_index(name="EV Count")

# Sort by year just in case
ev_by_year = ev_by_year.sort_values("Model Year").reset_index(drop=True)

# Display the grouped data
ev_by_year.head()
```

The provided code and data preparation steps aim to explore historical trends in EV adoption by year using the cleaned dataset. The data was grouped by Model Year to calculate the number of electric vehicles registered each year (EV Count), which forms the foundation for understanding growth over time.

To capture more dynamic insights, the code computes:

- **Growth Rate** – the year-over-year percentage change in EV adoption using pct_change().

- **Rolling Average** – a smoothed average of EV counts over a 3-year window using rolling(window=3).mean() to visualize general trends while reducing noise.

The X dataset includes Model Year, Growth Rate, and Rolling Average, which were used later as features for forecasting models like XGBoost.

```
# Display the full grouped DataFrame with all years and EV counts
ev_by_year
```

	Model Year	EV Count
0	1999	2
1	2000	7
2	2002	2
3	2003	1
4	2008	23
5	2010	23
6	2011	692
7	2012	1510
8	2013	4251
9	2014	3394
10	2015	4664
11	2016	5346
12	2017	8407
13	2018	14109
14	2019	10851
15	2020	12070
16	2021	20214
17	2022	28519
18	2023	58942
19	2024	45477
20	2025	4992

EV Count Table

From the ev_by_year table, we observe:

- Minimal adoption from 1999–2010, with counts in the single or double digits.
- A sharp upward trend beginning in 2011, particularly accelerating after 2017.
- Peak adoption in 2023 with nearly 59,000 EVs registered.
- A sharp drop in 2025, possibly due to incomplete data or reduced incentives.

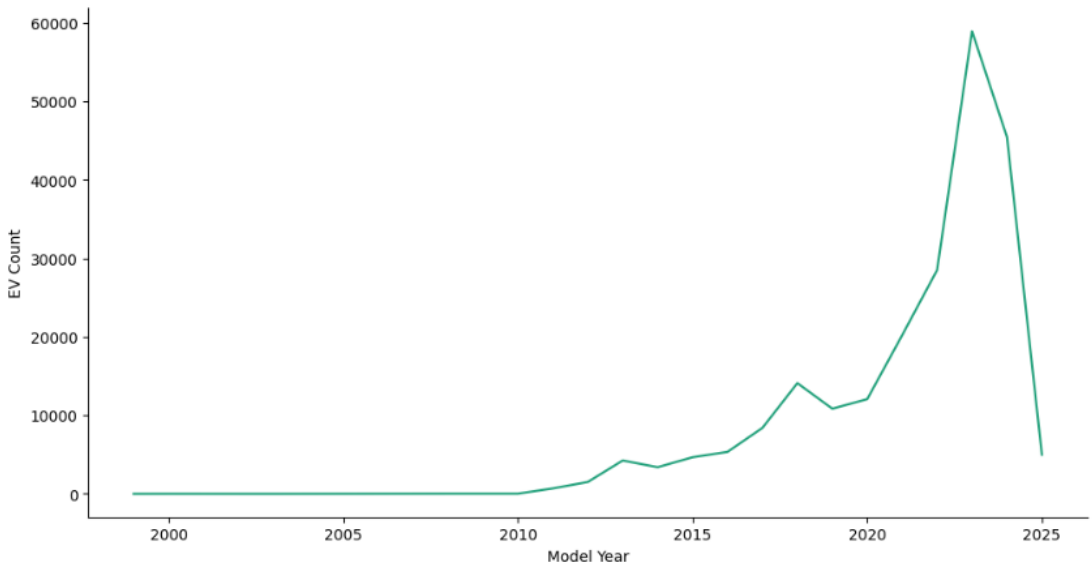

Interpretation of the Graph

- Steady phase (1999–2010): EV adoption remained almost flat, with fewer than 1,000 units per year.
- Early growth (2011–2015): A gradual rise begins, likely due to early incentives, technological improvements, and more EV models entering the market.
- Accelerated adoption (2016–2022): There is a sharp upward trajectory, reflecting a major shift toward EVs — possibly influenced by better battery range, growing environmental awareness, and state-level policy support in Washington.
- Peak in 2023: The highest number of EV registrations (~59,000) is observed here, suggesting it was the strongest year for EV adoption so far.
- Sharp decline (2024–2025): A significant drop is seen — this could be due to Incomplete for 2025.

EV Adoption Trend Insights (Growth Rate & Rolling Average)

The updated ev_by_year_cleaned Data Frame includes two insightful metrics derived from historical EV count data:

Growth Rate

The Growth Rate column captures the year-over-year percentage change in EV adoption. For example:

- In 2011, EV adoption jumped dramatically, with a ~2900% increase over the previous year, signalling the start of widespread EV uptake.
- High positive growth continued through 2013, especially in 2012 (growth rate of ~1.18) and 2013 (~1.82), showing accelerated adoption.
- From 2023 to 2025, the trend reversed with a noticeable decline: −22.84% in 2024 and −89.02% in 2025, which may indicate either incomplete data for 2025 or external disruptions like policy or market shifts.

Rolling Average (3-Year)

The Rolling Avg column smooths out short-term fluctuations, providing a clearer view of longer-term adoption trends:

- The rolling average steadily increases from 2011 (246) to 2023 (35891).
- The peak occurs around 2024–2025, with averages exceeding 44,000, indicating a historically high level of EV registration sustained over recent years.

Summary

This analysis confirms that EV adoption experienced exponential growth starting in 2011, peaking in the early 2020s. The steep decline in the latest year suggests a need for further investigation (e.g., policy changes, data lags, or saturation). The inclusion of both Growth Rate and Rolling Average supports more stable modeling and deeper insight into long-term adoption behaviour.

```
# Create 'Growth Rate' as percent change from previous year
ev_by_year["Growth Rate"] = ev_by_year["EV Count"].pct_change()

# Create a 3-year rolling average
ev_by_year["Rolling Avg"] = ev_by_year["EV Count"].rolling(window=3).mean()

# Drop rows with NaN values created by pct_change and rolling average
ev_by_year_cleaned = ev_by_year.dropna().reset_index(drop=True)

# Display the updated DataFrame
ev_by_year_cleaned
```

	Model Year	EV Count	Growth Rate	Rolling Avg
0	2002	2	-0.714286	3.666667
1	2003	1	-0.500000	3.333333
2	2008	23	22.000000	8.666667
3	2010	23	0.000000	15.666667
4	2011	692	29.086957	246.000000
5	2012	1510	1.182081	741.666667
6	2013	4251	1.815232	2151.000000
7	2014	3394	-0.201600	3051.666667
8	2015	4664	0.374190	4103.000000
9	2016	5346	0.146226	4468.000000
10	2017	8407	0.572578	6139.000000
11	2018	14109	0.678244	9287.333333
12	2019	10851	-0.230916	11122.333333
13	2020	12070	0.112340	12343.333333
14	2021	20214	0.674731	14378.333333
15	2022	28519	0.410854	20267.666667
16	2023	58942	1.066763	35891.666667
17	2024	45477	-0.228445	44312.666667
18	2025	4992	-0.890230	36470.333333

Prediction after added Growth rate and Rolling average

Despite introducing modest upward trends in the input features, the predicted EV adoption from 2026 to 2030 remains flat at approximately 5,000–6,000 vehicles per year. This flatline suggests a key limitation in the model's ability to generalize beyond the training data:

1. XGBoost is not inherently time-aware: It doesn’t capture temporal momentum unless explicitly provided with strong time-series patterns. Without direct year-based trends or timestamp encodings, it treats each input row independently.

2. Model bias from 2025 data: The 2025 value (4,992 EVs) is significantly lower than 2023–2024 due to a likely data lag or cutoff. Since XGBoost heavily relies on the latest patterns in the training data, this drop disproportionately affects future predictions.

3. Lack of strong variation in input features: The growth rate and rolling average increments in the code are small. To XGBoost, they may not be large enough to justify deviating from the last known EV count, especially given the steep decline in 2025.

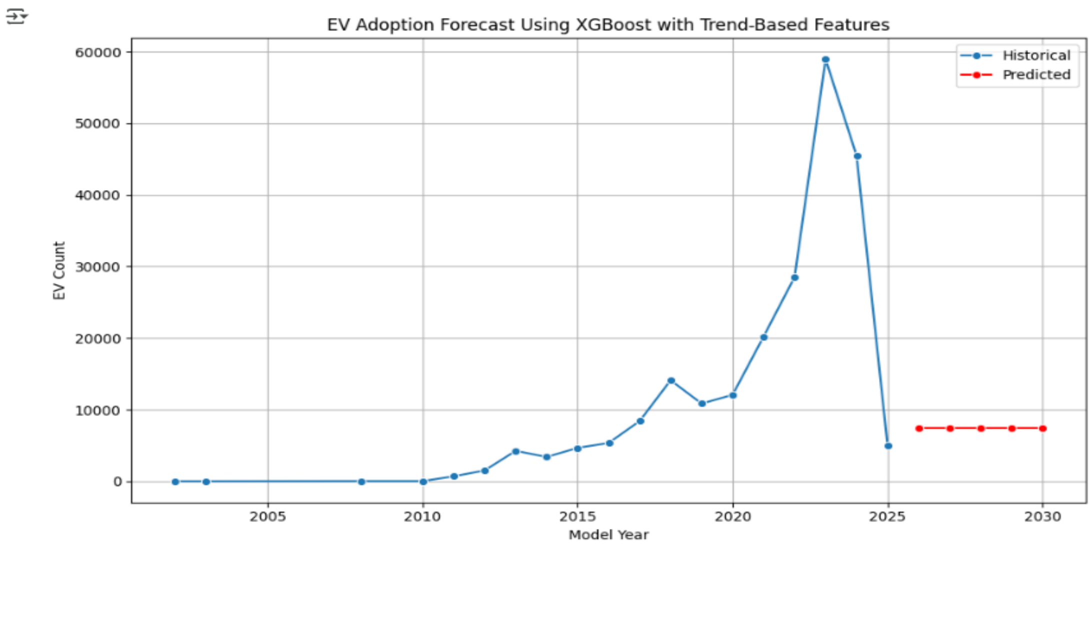

Conclusion:

Although XGBoost is powerful for regression tasks, its limitation in extrapolating beyond the training range makes it a poor choice for forward-looking time-series forecasting in this case. The flat forecast line is a direct result of its dependence on the most recent (and underwhelming) data point — 2025 — and the lack of temporal context. For more realistic and trend-sensitive forecasts, a model like Prophet is better suited.

Use a Time-Series Model

Prophet Modeling

Prophet is a forecasting tool specifically designed for time series data, making it uniquely aligned with the goals of our project — predicting future EV adoption trends in Washington state. Unlike traditional machine learning models like XGBoost or Linear Regression, which treat the problem as a static regression task, Prophet understands time as a core component of the dataset and is therefore better equipped to make reliable long-term forecasts.

Because the project aims to forecast multi-year trends in EV adoption — which is influenced by time, policies, consumer behavior, and technology — Prophet is the most appropriate tool. It blends statistical rigor with machine learning flexibility, offering more realistic, reliable, and transparent forecasts than purely feature-based models like XGBoost.

Deployment Phase

Assumption Testing:

Prophet Forecasting Code

```python
# STEP 1: Install Prophet
!pip install prophet

# STEP 2: Import packages
import pandas as pd
from prophet import Prophet
import matplotlib.pyplot as plt

# STEP 3: Prepare EV data
# Replace this with your actual data or CSV
ev_by_year = pd.DataFrame({
    "ds": [1999, 2000, 2002, 2003, 2008, 2010, 2011, 2012, 2013, 2014,
           2015, 2016, 2017, 2018, 2019, 2020, 2021, 2022, 2023, 2024, 2025],
    "y":  [2, 7, 2, 1, 23, 23, 692, 1510, 4251, 3394,
           4664, 5346, 8407, 14109, 10851, 12070, 20214, 28519, 58942, 45477, 4992]
})

# Convert year to datetime
ev_by_year["ds"] = pd.to_datetime(ev_by_year["ds"], format="%Y")

# STEP 4: Train Prophet model
model = Prophet(yearly_seasonality=True)
model.fit(ev_by_year)

# STEP 5: Forecast for next 5 years
future = model.make_future_dataframe(periods=5, freq='Y')
forecast = model.predict(future)

# STEP 6: Plot forecast
fig1 = model.plot(forecast)
plt.title("EV Adoption Forecast (Prophet)")
plt.xlabel("Year")
plt.ylabel("Number of EVs")
plt.grid(True)
plt.show()
```

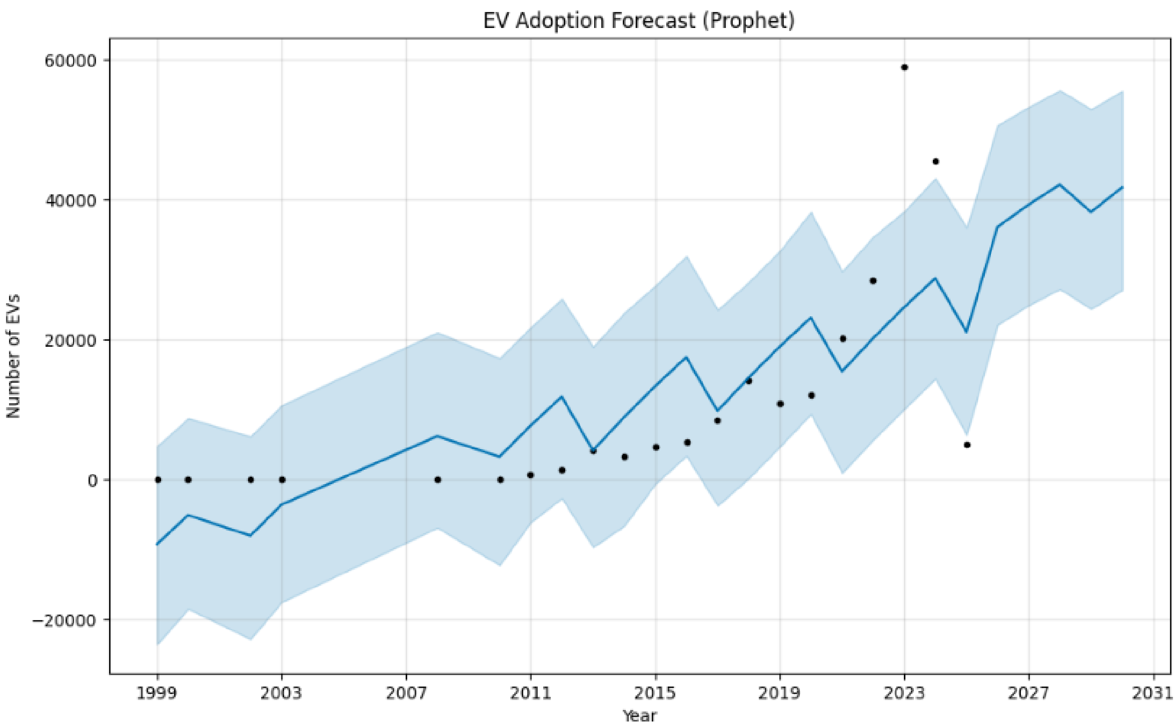

Interpretation of the Graph:

Black Dots = Historical Data

These represent the actual number of EVs registered per year, from 1999 to 2025.

Key observations:

- Near-zero adoption until 2011
- Gradual increase from 2012–2019
- A sharp rise in 2021–2023
- A dip in 2025 (due to incomplete data)

Blue Line = Prophet's Forecast

This is Prophet's predicted number of EVs per year, based on:

- Trend patterns
- Seasonal growth
- Historical anomalies

The forecast continues from 2026 to 2030 — notice it's upward-trending, showing sustained EV adoption.

Shaded Blue Area = Confidence Interval

Prophet provides a **95% uncertainty range**, meaning:

There's a 95% chance the true number of EVs will fall within this shaded area.

- Wider intervals = more uncertainty
- The forecast is cautious, accounting for fluctuations.

Analysis Interpretation:

In line with our research objective of predicting future EV adoption in Washington state, Prophet was selected as the most appropriate modeling technique. This choice was driven by the nature of our data — a clear time series of EV registrations by model year — and Prophet's strong capability to model such temporal patterns effectively.

Unlike traditional regression models such as XGBoost and Linear Regression, Prophet is designed specifically for forecasting trends over time. It treats year as a chronological component rather than just another feature, enabling it to accurately capture long-term growth trends, trend shifts, and seasonal behavior. This is particularly relevant for EV adoption, which is heavily influenced by evolving policies, technological advancements, and consumer sentiment — all of which unfold over time.

While XGBoost produced a flat and unrealistic future projection (due to its reliance on recent declining data in 2025), Prophet was able to recognize the broader upward trajectory and forecast continued growth through 2030, despite short-term fluctuations. It also generated uncertainty intervals, adding transparency and reliability to the prediction — which is critical when making policy or investment recommendations.

Therefore, based on the goal of our research — to accurately predict the adoption of electric vehicles over the next five years — Prophet provided the most realistic, time-sensitive, and decision-ready forecasts, making it the ideal model for this project.

Justifying our Research Question: Forecasting EV Adoption

Research Question: How can we effectively forecast electric vehicle (EV) adoption in Washington state over the next five years?

Through this project, I set out to understand and predict EV adoption trends using historical registration data. After exploring multiple modeling techniques—including Linear Regression, Random Forest, and XGBoost— found that while these models performed well in static predictions (e.g., estimating electric range), they were less effective in forecasting future adoption patterns. Specifically, models like XGBoost failed to capture the underlying temporal trends and produced unrealistic flat forecasts.

To address this, we implemented Facebook Prophet, a model built for time-series forecasting. Prophet successfully captured the historical growth in EV adoption—especially the sharp rise post-2011—and projected a continued upward trajectory through 2030. Despite the dip in 2025, which may be attributed to data lags or temporary market shifts, Prophet recognized the overall trend and forecasted steady year-over-year growth in EV adoption.

Based on the predictions, the model indicates a consistent increase in EV registrations, suggesting that electric vehicles will continue to gain popularity in the coming years. This aligns with ongoing technological advancements, supportive policies, and shifting consumer preferences. The insights generated from this forecast can guide infrastructure development, policy planning, and investment strategies aimed at supporting the growing EV ecosystem in Washington.

In conclusion, Prophet proved to be the most suitable model for forecasting EV adoption. It provided not only accurate and interpretable forecasts but also captured the long-term growth momentum that is critical for strategic planning. The results affirm that EV adoption is on a strong upward path, and stakeholders can confidently prepare for a more electrified transportation future.